WORKING GUIDE

The Pyramid Principle

A working guide to structuring arguments for senior audiences

---

Barbara Minto's method · An applied guide with examples from an analytics function in financial services · Updated November 2026

Contents

0. How to use this guide

PART I — THE METHOD

2. Why the pyramid: the underlying argument

3. Anatomy of the pyramid

4. MECE: the discipline of grouping

5. The introduction: the SCQA pattern

6. Building the pyramid: two routes

7. Ordering ideas within a group

8. Summarising groups: where most documents are lost

9. Deduction and induction

10. Before structuring: defining the problem properly

11. Getting the pyramid onto the page

12. The pyramid in presentations

13. Adapting to a regulated financial environment

PART II — APPLIED CASES

14. Three worked examples

15. Data quality when someone else owns the data

16. Getting a decision in a globally centralised organisation

17. Building a business case for funding (funding request)

18. Difficult conversation downward: asking someone on your team to change a behaviour

19. Difficult conversation upward: when the problematic behaviour is your manager's

20. Structuring a "no" (saying no)

21. Handing over a project or a system (handover, knowledge transfer)

22. Explaining a technical subject in plain terms (pitching to a non-technical audience)

PART III — PRACTICE

23. Common failures in a quantitative profile

24. Checklist

25. A 30-day practice plan

26. Running this as a team workshop

27. What this method does not solve

28. Glossary

29. Sources and attribution

Barbara Minto's method · An applied guide with examples from an analytics function in financial services · August 2026

0

How to use this guide

This is not a summary of Barbara Minto's book. It is an independent guide to applying the method: it explains the ideas in its own words, with examples built around real situations in an analytics function inside a financial firm — risk committees, memos, model notes tied to credit risk measurement but generalisable to other uses, escalations, difficult conversations — because the original material uses corporate examples from another era that make the transfer harder than it needs to be.

It is written to be read alone and also to be worked through as a team. Part I is the method. Part II applies it to eight recurring situations in any function that depends on others for its data and on a head office, often in another geography, to approve its decisions — which means aligning the group strategy with the subsidiary's. Part III is practice, and includes a workshop agenda for anyone planning to teach the method to their team.

The method is not learned by reading. It is learned by rewriting. Section 24 is a one-page checklist to keep open while you write; section 25 is a 30-day practice plan. If you only use two parts of this guide, use those.

> **The rule that summarises the whole guide: if your reader has to reach the fourth paragraph to find out what you are asking for, you have already lost. The conclusion goes first; everything else exists to support it.**

PART I

The method

How to build an argument that lands the first time.

2

Why the pyramid: the underlying argument

2.1 If the text imposes no structure, the reader imposes their own

The human mind cannot hold many loose items at once. Faced with nine findings and no hierarchy, a reader starts inventing categories in order to retain them: by topic, by order of appearance, by whatever worries them most. That improvised grouping almost never matches yours.

Hence the central claim: if you do not impose a structure, the reader imposes a worse one. The pyramid is not a stylistic preference. It moves the cost of ordering from the reader to the writer, who is the one with the context to do it well.

2.2 People read to answer a question

Nobody reads a memo out of curiosity. They read it because they have an open question — do I approve this, should I worry, what do I do — and they are looking for the answer. A well-structured document reads like a dialogue: the idea above raises a question in the reader's mind, and the level immediately below answers it.

When that dialogue works, the reader can stop at any depth and still hold a complete, coherent version of the argument. A director reads three lines and can decide. An analyst goes three levels down and finds the detail. The same document serves both.

2.3 What a quantitative profile gains specifically

The typical failure of a technical analyst is not a lack of rigour. It is reproducing, in the document, the order in which the work was done: data, then methodology, then results, with the conclusion at the end — the order of a paper. That order is right for proving something and poorly suited to deciding something.

The pyramid inverts the sequence without sacrificing rigour. The conclusion leads; the rigour is still there, but organised as support for the conclusion rather than as the road towards it. In practice the difference is that, in a twenty-minute committee, there is time to say the thing that matters.

3

Anatomy of the pyramid

A well-built pyramid satisfies three rules at once. Each is verified differently, so it helps to keep them apart. The single idea at the apex is called the *governing thought*, or *main message*: the statement that governs everything below it.

3.1 The vertical rule: every idea summarises the ones below it

An idea at any level must be the summary of the ideas immediately below. Not their title, not their category: their summary. If you deleted the level below entirely, the idea above should still communicate the message.

> *Fails: "We identified three problems in the portfolio."*
>
> *Works: "The deterioration is concentrated in a single sector and is reversible through limit adjustments."*

The first version forces you to read three bullets before you know anything. The second states the message; the three bullets now support it, which is their actual job.

3.2 The horizontal rule: ideas at the same level are of the same kind

Items hanging from the same idea must belong to one logical category: all reasons, or all steps, or all problems, or all criteria. Mixing a reason with a step and a recommendation produces a list that feels disordered even when every item is true.

Quick test: try to give the group a plural noun. "Three reasons", "four steps", "two constraints". If the only noun that fits is "things" or "points", the group is not homogeneous.

3.3 The order rule: every group follows an explicit logic

Within a group the order cannot be arbitrary. There are only three legitimate orders, and each one reveals the kind of analysis behind it. Section 7 develops them.

4

MECE: the discipline of grouping

MECE stands for Mutually Exclusive, Collectively Exhaustive. A group of ideas is MECE when the items do not overlap and, together, they cover the whole relevant space.

4.1 The two halves, separately

- Mutually exclusive: no item contains or repeats part of another. If explaining point two means going back over point one, they overlap.
- Collectively exhaustive: nothing the reader would consider part of the problem is left out. The honest test is to ask what the best-informed sceptic in the room would say is missing.

4.2 Cuts that meet the MECE standard

Robust groupings usually come from a pre-existing, verifiable structure rather than from brainstorming:

- An algebraic identity: expected loss decomposes into probability of default, loss given default and exposure. MECE by construction.
- A partition of the population: by segment, vintage, geography or rating band. MECE if the categories do not overlap and add to the total.
- A process sequence: origination, monitoring, collections. MECE if the process is well defined.
- A dichotomy: internal versus external, controllable versus not, short term versus long term.

4.3 The three usual failures

1. The "other" bucket. It almost always means the cut was badly chosen. If "other" is more than ten per cent, the cut should be redone.
2. Hidden overlap. "Data problems", "system problems" and "governance problems" look clean, but a field mis-captured by a system with no owner falls into all three.
3. False exhaustiveness. You cover everything you analysed, not everything that exists. A senior reader spots the gap in seconds and distrusts the rest from that point on.

> **An honest caveat: MECE is a verification standard, not a publication requirement. Forcing an artificial MECE grouping produces rigid, unnatural documents. Use it to test a grouping that already makes sense, not as a template to generate one.**

5

The introduction: the SCQA pattern

SCQA stands for Situation, Complication, Question, Answer. It is the mechanism that installs, in the reader's mind, the question your document answers. Without it, an opening conclusion sounds abrupt or presumptuous.

5.1 The four elements

SITUATION

A fact about the subject that the reader accepts without argument. Not general background, not history: the point of agreement. It must be uncontroversial enough that the reader nods along. If they can object to the situation, you are not building a base, you are opening a front.

> *"The credit card origination model was approved in 2021 and now decides ninety per cent of digital-channel applications."*

COMPLICATION

What disturbed that stable situation and why we are having this conversation. It is the engine of the document: a regulatory change, an observed deterioration, a deadline, a validation finding, a pending decision.

> *"Extending it to the payroll-lending product, due to launch in March, requires recalibrating over a window that includes 2020, and internal validation found that the current calibration understates default in the intermediate risk bands."*

QUESTION

The logical consequence of the complication. In a well-built document it is not written: it arises on its own in the reader's mind. Writing it out is optional and sometimes useful in long documents, but if you have to force it, the complication is not the right one.

> *"Do we recalibrate, rebuild, or apply a temporary conservative adjustment?"*

ANSWER

The conclusion, the *Answer* of the pattern. It is the idea that leads the pyramid — the *governing thought* — and the first substantive paragraph of the document. It must be a statement with content, not a promise of content.

> *"We recommend recalibrating with a transitional conservative adjustment: it preserves the launch date, costs a fraction of a rebuild, and the impact on provisions is bounded."*

5.2 The four openings and when to use each

S–C–Q–A is the standard order, but it can be rearranged depending on what you know about the reader. The content does not change; the entry point does.

Opening	Order	When it works
Standard	Situation → Complication → Answer	A neutral audience that needs context. The default case.
Direct	Answer → Situation → Complication	The reader knows the problem and is short of time. Crowded committee agendas.
Concerned	Complication → Situation → Answer	There is real urgency and you want it felt. Deterioration, incident, date at risk.
Sceptical	Question → Situation → Complication → Answer	The reader doubts the framing itself. Naming the question disarms resistance.

5.3 Introduction failures

Four recurring failures:

- **The résumé opening:** setting out all the work done before saying what was found. The reader does not need an account of the effort in order to judge the finding.
- **An over-long situation:** three paragraphs of background before the complication. The situation should take two or three lines.
- **A missing complication:** if what changed cannot be named, the document was probably not needed.
- **A deferred answer:** "conclusions are presented in section 5". That postpones the answer and forces the reader through the whole document to find it.

6

Building the pyramid: two routes

There are two ways to arrive at the structure, *top-down* and *bottom-up*. They are used at different moments and both are legitimate.

6.1 *Top-down* (when the conclusion is already known)

Six steps:

1. Write the subject of the document in one line.
2. Write the question the reader has about that subject.
3. Write your answer to that question, even provisionally.
4. Check the introduction: can you state a situation the reader accepts and a complication that generates exactly that question? If not, your answer does not address what they want to know.
5. Anticipate the next question your answer raises — usually "why?" or "how?" — and write the three or four points that answer it.
6. Repeat down each branch until the next question no longer interests the reader.

6.2 *Bottom-up* (when there are loose findings)

Six steps:

1. List every point you want to communicate, unedited, one per line.
2. Identify the relationships: which cause which, which are instances of the same phenomenon, which are consequences.
3. Group the ones that share a nature and give each group a plural noun: reasons, symptoms, options, constraints.
4. For each group write one sentence summarising what they say together — not what they are, what they say. That sentence moves up a level.
5. Repeat until you reach a single idea. That is your answer.
6. Now build the introduction backwards and check that the question it generates is the one your idea answers. If they do not match, the grouping broke somewhere and needs redoing.

In practice the two are used together: you come down from a hypothesis, hit evidence that does not fit, work back up from the findings, and adjust the conclusion. What matters is closing the loop at least once before publishing.

7

Ordering ideas within a group

There are only three logical orders. Any other reproduces the sequence in which the ideas happened to arrive. Choosing one of the three deliberately is also a way of auditing your own analysis: the order reveals the kind of reasoning behind it.

7.1 Time order (chronological or causal)

The items are steps that happen one after another, or linked causes. Used when describing a mechanism, a flow or an implementation sequence.

Validity test: if two items are swapped, does the argument stop making sense? If they admit any order, it was never a process: it was a list presented as a sequence.

7.2 Structural order (parts of a whole)

The items are components of a structure: regions, business lines, terms in a formula, layers of a system. Used when decomposing something that exists.

Validity test: do the parts add up to the whole without overlap? If they do not, the cut is worth correcting before going further, because the flaw travels up to the levels above.

7.3 Degree order (by importance)

The items share a property and are ordered by how much of it they have: most severe to least, largest impact to smallest. Used when ranking or prioritising.

Validity test: is there a metric, even a qualitative one, that justifies the order? If it cannot be named, the order rests on an impression rather than a criterion. The metric almost always exists: size, cost, impact, urgency or likelihood of occurrence. In a risk function, for instance, the natural order is usually by exposure or by capital impact.

> **Rule of thumb: if a group admits none of the three orders, it is worth rethinking the grouping before writing any more. Redoing it now costs less than discovering the gap in the committee room.**

8

Summarising groups: where most documents are lost

This is where most people abandon the method without noticing. The pyramid gets built, the grouping is sound, and then each group is labelled with an empty tag. The result keeps the shape of a pyramid, but the reader still does not receive the message.

8.1 Intellectually blank assertions

The term is Minto's. It names the sentences that announce the existence of content without delivering it, and "blank" refers to the sentence — it carries no idea — not to whoever wrote it. They are recognisable because the reader cannot do anything with them.

Blank	With content
The company has three objectives.	The three objectives are mutually incompatible at the current budget.
We analysed model performance.	The model discriminates well but calibrates badly: the ranking holds, the level does not.
There are several regulatory considerations.	Neither route meets the regulatory deadline without a transitional adjustment.
We identified improvement opportunities.	Two controls have had no formal owner since the reorganisation.

8.2 The two summary rules

- Action groups (steps, recommendations, activities): the idea above must state the result of carrying out all the steps. Not "four actions for the project", but "these four actions close the project before the regulatory date".
- Situation groups (findings, facts, evidence): the idea above must state the inference that follows from all of them. Not "three validation findings", but "all three findings point to the same origin: the calibration window".

8.3 The "So what?" test

After writing any summary idea, ask it out loud: so what? If you have an answer, that answer was the idea you should have written. Repeat until the question stops producing anything new.

> *"Mid-segment arrears rose 40 basis points." → So what?*
>
> *"It is the third consecutive quarter of increases." → So what?*
>
> *"The deterioration is structural, not seasonal, and the current cut-off no longer protects the target return." ← This is the idea.*

This is the highest-return exercise in the guide and also the most uncomfortable, because it shows up the times when there is no conclusion yet, only numbers. When the exercise produces nothing new, the cause usually lies in the analysis rather than the writing: it is worth going back to the data before writing any further.

> **A conclusion without a period, a universe and a scope cannot be verified. "Arrears went up" cannot be discussed; "mid-segment arrears rose 40 basis points in the third quarter" can.**

9

Deduction and induction

A group of ideas holds together in one of two ways, and it is worth knowing which one you are using, because they summarise differently.

9.1 Deduction

A chain: a statement, a second statement about the first, and the conclusion that follows from both. Each step depends on the one before.

> *The calibration window does not cover a full cycle. → A model calibrated without a full cycle understates default in a downturn. → The current calibration understates default in a downturn.*

It is rigorous and very convincing in detail. The cost is that the reader must hold the whole chain in mind before receiving the message, and the argument breaks entirely if one link fails.

9.2 Induction

A set of items of the same kind from which a common conclusion follows. There is no chain: each item supports the conclusion on its own.

> *The window does not cover a cycle. The sample excludes the mid segment. The cut-off has not been reviewed in three years. → All three share one origin: the calibration is anchored to conditions that no longer hold.*

It is easier to follow, more robust — if one item falls, the conclusion survives — and it lets you state the conclusion first.

9.3 Which to use

Rule of thumb: induction at the top levels, deduction at the lower ones. Senior readers look for the pattern; technical reviewers look for the chain. A document that uses deduction at the conclusion forces everyone to read all of it, which is precisely what the method sets out to avoid.

10

Before structuring: defining the problem properly

A good share of badly structured documents did not fail at the structuring stage: they failed earlier, at defining the problem. They are well built and answer precisely, but they answer a different question from the one the reader has. This section is the antidote.

10.1 The definition sequence

The five questions serve at two moments. **Before writing**, to define the problem: ten minutes spent here saves entire rewrites. And **on a draft that is not quite working** — your own, or one somebody asked you to review — as a diagnostic: the question that cannot be answered usually points to exactly where the fault lies.

1. Starting point: what was the stable situation, and who is the actor?
2. Disturbing event (*disturbing event*): what changed?
3. Undesired result (R1): where are we today as a consequence?
4. Desired result (R2): where do we want to be, on what metric and by when?
5. Question: what exactly does the reader need to decide or know?

The distance between R1 and R2 is the problem. If that distance cannot be measured, the document has no reference point. The desired result has to be quantifiable — a metric, a target value and a date; without those three there is no way to know whether the problem was solved.

The question in point 5 determines how the work proceeds. It almost always falls into one of two types: *why is this happening*, which calls for a diagnosis, or *what do we do*, which calls for comparing options. Each type is worked with a different tool, and that is what the next two sections develop. If that question is still unclear, it is worth settling first: neither tool compensates for a badly defined problem.

10.2 Diagnostic frameworks

When the question is "why is this happening", decompose the problem into a tree before investigating. Three ways to build one:

- By mechanism: reproduce how the process or system actually works and locate where it breaks. The most powerful and the most demanding of business knowledge.
- By possible cause: enumerate the categories of cause that could produce the observed effect, then rule them out with evidence. Useful when the mechanism is unclear.
- By undesirable elements: list what should not be happening and trace each to its origin. Useful in incidents.

10.3 Logic trees for decisions

When the question is "what do we do", the tree decomposes options rather than causes, with explicit evaluation criteria. Its advantage over a narrative analysis is that it forces the criterion to be declared before the outcome is known, which is the only way to compare alternatives without bias.

11

Getting the pyramid onto the page

The mental structure has to become visible. If the page does not show it, the reader does not receive it, however good the underlying idea. Four devices do that work: headings, the executive summary, transitions between blocks, and visual hierarchy.

11.1 Headings that are ideas, not topics (*message titles*)

A heading should stand on its own and communicate a complete idea. "Methodology" is a topic. "The recalibration uses a window extended to twelve years" is an idea. If the reader scans only the headings, they should come away with the whole argument. That is the standard.

The consequence: a twenty-page document should have eight to twelve headings with real content, not forty generic labels.

> ***What the improvement looks like.*** *The headings of a note about a calculation error, first as topics:*
>
> *Background · Methodology · Findings · Impact · Recommendations*
>
> *And the same ones, turned into ideas:*
>
> *1. We recommend fixing the calculation at source before the annual close*
>
> *2. The error originated in the 2022 integration and affects three of the twelve reports*
>
> *3. Fixing it at source costs a third of what rebuilding the platform would cost*
>
> *4. The fix fits the calendar without moving the close date*
>
> *The first list says nothing on its own. The second, read straight through and without the body of the document, is already the complete argument. That is the test.*

11.2 Executive summary (*executive summary*)

Not an abbreviated version of the document, and not a list of sections. It is the conclusion plus the first level of support, in prose. Four to eight lines. If the reader reads only that, they should be able to take the decision and explain it to someone else.

A sufficiency test: hand only the executive summary to someone unfamiliar with the subject and ask them what decision is being requested and why. If they hesitate, the summary needs rewriting.

11.3 Transitions

Moving from one block to the next, the reader needs to know where they stand in the argument. One sentence is enough: pick up the idea just closed and announce the question the next one opens. What does not work is the decorative transition, which takes up space without orienting anyone ("The following section presents...").

11.4 Numbering and visual hierarchy

Three rules:

- Visual hierarchy must match logical hierarchy. Two ideas at the same level look the same; a subordinate idea looks subordinate.
- It is best not to go past three levels of depth in a business document. If a fourth is needed, that material belongs in an annex.
- Annexes exist to move out of the body what supports but does not argue: result tables, technical specification, robustness checks. Their existence does not remove the need to summarise their message in the body.

12

The pyramid in presentations

Adapting to slides is where the method pays off most visibly, because a committee does not read: it listens and looks. The next three sections cover the slide title, the one-idea-per-slide limit, and how to handle charts.

12.1 The slide title is the message (*action title*)

The rule, known as the *action title*, is simple and seldom applied: the title of each slide is that slide's conclusion, written as a full sentence. The body of the slide is the evidence for the title.

> *Weak title: "Arrears evolution by vintage"*
>
> *Correct title: "Vintages originated after the policy change deteriorate 30% faster"*

The practical consequence: reading only the titles of the presentation aloud, in order, should produce a complete and coherent argument. That is what is called the *storyline*, and it is the best diagnostic available for the material as a whole. Worth doing before every committee; it takes ninety seconds.

12.2 One idea per slide

If the title needs a conjunction to be true — "and also" — it is two slides. The temptation to compress comes from wanting to look efficient, and the effect is the opposite: each idea needs its own evidence, and one slide cannot hold both.

> ***A title that calls for splitting:***
>
> *"The slowdown is concentrated in discretionary consumption and the annual forecast must be revised before September"*
>
> *These are two ideas with different evidence: the first is supported by the breakdown by sector, the second by the impact on the forecast. Together on one slide, one of the two is left without support on screen. Separated:*
>
> *1. The 2024 slowdown is concentrated in sectors tied to discretionary consumption*
>
> *2. The annual forecast needs revising before the September publication*

12.3 Charts

Three rules:

- The chart title states the finding, not the variable plotted.
- Highlight visually the element that should be looked at. If everything has equal visual weight, the reader decides where to put their attention.

- A chart that needs spoken explanation does not work in a document that will circulate without its author, which is what happens to almost all material in a large organisation.

13

Adapting to a regulated financial environment

The method was designed for consulting. Transplanting it unchanged into a supervised financial institution creates real friction worth anticipating: the reader may be a supervisor who expects an established order, the conclusion may be uncomfortable to state in the first line, and the same document may have to serve two reporting lines with different interests.

13.1 When the reader is the regulator

A supervisor does not read to decide quickly: they read to verify compliance and traceability, and they often expect the document to follow an established order. In that case the sequence cannot be freely rearranged.

The solution is not to abandon the pyramid but to double the entry: a pyramid-format summary at the front, and the body in the order the supervisor expects. The summary provides clarity; the body, compliance. In practice, a conclusion at the front rarely draws an objection: what the supervisor requires is that the body keep the expected order.

13.2 When the conclusion is unwelcome

The most common objection to "conclusion first" is political: if the message is bad, saying it in line one feels blunt. The objection is legitimate and the answer is that the order does not change, the framing does — the conclusion still goes first, but stated together with the way out.

> *Blunt: "The model will not pass validation."*
>
> *Pyramid with framing: "The model will not pass validation in its current form; a transitional conservative adjustment preserves the regulatory date while it is recalibrated."*

Both say the same thing. The second is just as honest and leaves the reader holding a decision rather than a problem to solve.

13.3 Dual reporting lines

When the same matter goes to two reporting lines with different interests, the temptation is to write something neutral that unsettles neither. The usual result is a document so general that neither of them finds in it what they need.

The pyramid offers a way through: the conclusion — the fact and the recommendation — is the same for both; what changes is the introduction, because the relevant complication differs for each. One document, two entrances. This is not a matter of holding two positions, but of recognising that each reader arrives with a different question.

PART II

Applied cases

Eight situations that recur in any function depending on others for its data and on a head office to approve its decisions. The examples are built in a financial setting, but the structure of each case transfers to other sectors unchanged. Each one covers why the intuitive approach fails, how it is structured, a worked example and the typical errors.

14

Three worked examples

14.1 Sector analysis memo

Typical version, written in work order:

> BEFORE
>
> *"The quarterly review of sector activity was performed. 1,240 companies were reviewed across the ten sectors contributing most to gross domestic product (GDP), with data as of June. The methodology consisted of rates of change by industry and analysis by company size. Results show variations across several sectors. Conclusions and recommendations are presented in section 4."*

Pyramid version:

> AFTER
>
> *"The slowdown in the **third quarter of 2024** is concentrated in a single sector and reflects a temporary factor.*
>
> ***Across the 1,240 companies reviewed — the ten sectors contributing most to gross domestic product (GDP) —**, export manufacturing accounts for 70% of the fall, while the rest remains stable. The origin is a supply-chain interruption, not a contraction in domestic demand. The 12 largest companies in the sector account for half the effect.*
>
> *A decision is needed on whether to revise the annual forecast before the 15 September publication."*

The order changed, not the content. The first version demands twenty minutes of reading to arrive at the same place; the second leaves the reader able to decide in thirty seconds and to go deeper if needed. The framing is split: the period sits in the conclusion, the universe in the first line of support.

14.2 Data platform proposal note

SCQA introduction:

- *Situation:* the reporting platform was implemented five years ago and feeds the twelve closing reports.
- *Complication:* the annual review found that the balance calculation has been wrong since the 2022 integration and affects three of those reports; the closing calendar does not allow the whole platform to be rebuilt.
- *Question (implicit):* fix at source, rebuild the platform, or apply a manual adjustment?
- *Answer:* fix at source, with a transitional manual control.

First level of support — three reasons, in degree order, heaviest first:

1. It is the only one of the three routes that fits before the annual close.

2. The diagnosis points to the integration, not the design of the platform: rebuilding it would attack a problem we do not have.
3. The cost is a third of a rebuild, and the manual control is withdrawn as soon as the fix is in production.

Note that the three reasons are of the same kind, do not overlap, and together cover the three dimensions a committee evaluates: time, diagnosis and cost. That is MECE applied to a grouping that already made sense.

14.3 A short email to a functional manager

The method works just as well in five lines. A well-structured decision email:

> *"Please could you confirm your decision on the data contract renewal before Friday.*
>
> *The supplier confirmed the price rises 18% if we renew in December and 32% if we wait until March. That leaves two options: renew now for two years at the current price, or go to tender and absorb the increase.*
>
> *My recommendation is the first: the saving is verifiable and the tender timeline does not fit before expiry. If you prefer the second, I need to open the process this week.*
>
> ***Happy to talk either option through.**"*

Courtesy does not compete with structure: an opening "please could you" and a brief close both fit without moving the request from its place. What does not work is courtesy taking up the whole first paragraph, with the request appearing in the third.

The structure is the same as in a long document: the request leads, context follows, the recommendation is explicit and each alternative carries its consequence. What changes is the length, not the method.

15

Data quality when someone else owns the data

This is probably the most frequent and worst-handled situation in an analytics function: the work depends on information produced by another team, the information arrives flawed, and there is no formal authority over whoever generates it. The usual outcome is an email listing defects that the owning team receives as a reproach, answers defensively, and the problem remains. The next five sections cover why the intuitive approach fails, how to structure the request, which rules improve the response rate, how to escalate without damaging the relationship, and a worked example.

15.1 Why the intuitive approach fails

The intuitive approach puts the defect in the conclusion: "we found 43 inconsistencies in the collateral file". Framed that way, the recipient is placed in the position of the accused and the writer in that of an uninvited auditor. The message also contains no request: the reader does not know what is expected of them, and the email gets filed.

The correct conclusion is not the defect, but the decision or deliverable at risk, together with the bounded action that unblocks it. The defect drops to the second level, where it belongs: it works as evidence, not as the message.

15.2 Proposed structure

- **Situation:** name the shared process in terms the other party recognises as their own. *"The monthly collateral file feeds the loss-given-default calculation — the estimated loss in the event of default — and the September regulatory report."*
- **Complication:** the measured defect plus its dated consequence. Measured means a rate and a volume, not adjectives. *"Twelve per cent of records — 340 out of 2,800 — arrive with no valuation date; at that level the calculation cannot be signed off before 15 September."*
- **Answer:** the smallest request that solves the problem, addressed to one person, with a date. Not *"improve the quality of the file"*, but *"complete the valuation date on the 340 records in the list I am sending you, before 5 September"*.
- **Support:** what has already been done on the requesting side. This is not self-promotion: it is what stops the other party feeling someone else's work is being handed over. *"We applied a temporary control that imputes the date by segment average; it works for internal monitoring, not for the regulatory calculation."*

15.3 Practical rules that can change the outcome

Six rules that tend to improve the response rate:

1. **Always quantify.** A defect with no rate and no impact registers as an impression, and an impression does not compete with the priorities the owning team has already committed to.
2. **Translate the defect into the other party's consequence.** What it costs the requesting side weighs little on their list; what compromises something they sign weighs a great deal. *"We have to reprocess"* rarely moves a priority; *"the report you sign would go out qualified — a formal audit observation —"* does.
3. **One request per counterpart.** An email with nine requests usually ends with none of them actioned: the recipient does not know where to start and postpones the whole thing.
4. **Prioritise, and make it visible.** With 43 defects, it helps to group them by origin — capture, system, definition — and flag which block the deliverable and which do not. Often three account for most of the impact.
5. **Offer the path of least effort.** Where it is possible to hand over the exact list of records to correct rather than asking for them to be found, it is worth doing. Half an hour of your own time costs less than two weeks of back and forth.
6. **Leave a record without an accusatory tone.** Documentation is necessary and needs no adjectives: dated facts are enough, and they are what supports the case if the matter escalates.

15.4 Escalation works in three levels and without surprises (*escalation path*)

Escalating is a normal coordination mechanism when two functions hold different priorities, not a sanction or a breach of trust. What damages the relationship is not escalating, but doing it without notice. The sequence that preserves it has three levels, and at every one the counterpart knows in advance what comes next:

1. **A direct request**, with a date and a consequence, and nothing else.
2. **A reminder** with the impact updated and an explicit note that, absent a reply, the matter will go to the relevant forum. Framed as information about the process, not as a warning.
3. **Escalation with options.** It is worth bringing two routes with their costs, not just the problem: an escalation that only reports it transfers the work to whoever decides; one that proposes alternatives hands them a decision.

15.5 Worked example

Typical version:

> BEFORE
>
> *"Good morning. Attached is the result of the quality review of the August collateral file. 43 issues of various types were detected and are detailed in the annex. It is worth mentioning that these issues have recurred in previous months. I look forward to your comments."*

Pyramid version:

> AFTER

> *"Please could you send the valuation date for 340 collateral items before 5 September, so that the loss-given-default calculation for the regulatory close can be signed off.*
>
> *The August file arrived with that field empty on 12% of records — 340 out of 2,800. We applied a segment-level imputation that works for internal monitoring but is not admissible for the reported calculation.*
>
> *The annex lists the exact 340 records, sorted by exposure. The first 60 account for 70% of the impact; if time is short, we can proceed with those.*
>
> *The remaining issues this month do not block the close; I will send them separately next week so they can be reviewed properly at source.*
>
> ***Many thanks.**"*

> **What changed: the conclusion moved from the defect to the decision at risk, the request became bounded and dated, the search work is already done, a partial route is offered if time runs short, and the urgent was separated from the important. The content is the same; what changes is that the recipient knows exactly what to do and by when.**

16

Getting a decision in a globally centralised organisation

When the decision is not taken in the local entity, the problem stops being technical and becomes one of competing for attention. Whoever decides at head office covers twenty entities, and the request competes with the other nineteen. The pyramid helps because in that context you have one paragraph before the proposal is filed into a category. The next four sections cover how to frame the conclusion, what the case must contain, who to address it to, and in what format.

16.1 The conclusion has to be framed as a problem the centre recognises as its own

The intuitive conclusion is "the local entity needs X". From head office, that framing reads as a local exception: cost to the group, benefit to a single entity — the category that wins the fewest approvals.

The framing that works presents the same request as the answer to a problem the group already recognises as its own. This is not about altering the substance, but about choosing — among several equally true descriptions of the request — the one that is relevant to whoever decides.

> *Local: "We need an exception to the model standard for the local portfolio."*
>
> *Corporate: "The global standard does not contemplate markets with limited history; this entity can be the pilot that closes that gap before three others raise it."*

16.2 The five elements of a case addressed to head office

1. **An anchor in a stated priority.** Cite the group objective, programme or policy the proposal aligns with. Without that anchor, the proposal competes with no institutional backing.
2. **Translation into the group's metric.** What is measured locally in local currency and operational load is assessed at head office in capital, in regulatory risk or in consistency across entities. The translation is part of the work.
3. **The smallest possible decision.** Rather than a new policy, it is better to request a pilot, a bounded exception with a review date, or a temporary waiver. What is reversible meets far less resistance.
4. **Handling the precedent.** Whoever decides at group level always assesses what happens if other entities ask for the same thing. It is worth anticipating that and defining the criterion that turns the case into a category rather than a concession: *"applicable to entities with less than five years of history in the segment"* is a criterion; *"applicable to this entity"* is a concession.
5. **The cost of not deciding, in the group's metric.** In a financial group, the argument that usually carries most weight is local regulatory risk, because head office answers for it. It is worth using where it applies and not overstating it where it does not: credibility lost by exaggerating once is hard to recover.

16.3 Identifying the three recipients

Before drafting, it is worth answering three questions in writing: who signs, who can veto without signing, and who is affected if the proposal is approved. Most proposals stall on the second group, which was not consulted and finds out in the forum.

Where dual reporting exists — a functional line into head office and an entity line in-country — attending to only one of the two is a frequent and costly error. The functional line provides technical feasibility; the entity line, institutional backing. Both are necessary, and both should see the version that concerns them before the forum, not during it.

16.4 One page: the request, three reasons, quantified impact and a date

The format that arrives intact is one page, with the request and the decision sought in the first paragraph, three reasons, the quantified impact, and the date after which the decision stops being useful. Everything else goes into annexes. A proposal that needs fifteen slides to be understood is not yet ready to go to the group; it is still internal working material.

17

Building a business case for funding (funding request)

A business case is a document of persuasion in the format of a technical document. Whoever assesses it has reviewed many before and readily recognises overstated benefits. The pyramid structure helps, but here credibility weighs as much as order. The next three sections cover how to open the document, what it must contain, and what sustains its credibility.

17.1 The document opens with the decision sought and its amount

The first paragraph answers four questions: what is being approved, what it costs, what it delivers and by when. A business case that opens with project background forces the reader to hunt for the request, and that hunt rarely happens.

> *"Approval is requested for a budget of 2.4 million to automate the regulatory reconciliation process over twelve months. The automation frees the equivalent of two manual positions, closes the audit finding open since March, and pays back in under two years."*

17.2 Proposed structure

Six elements, in this order:

1. **The current cost, quantified.** Not the problem in the abstract, but what doing nothing costs, in hours, in losses or in open risk. It is the number that holds up the whole document, and the one most often omitted.
2. **The options considered, including doing nothing.** A business case with a single option reads as a decision already taken in search of approval; three options with explicit criteria read as analysis.
3. **The recommended option, and why it is superior under those criteria.**
4. **Cost and benefit with the assumptions on view.** Every assumption behind the benefit, stated and available to be challenged. When assumptions are not visible, the assessor assumes they are optimistic.
5. **The consequence of not approving**, stated in concrete terms and without dramatising.
6. **The request and the real decision deadline.**

17.3 Credibility rules

Five rules that hold the document up in front of an experienced assessor:

- **State the benefit in the metric of whoever decides.** Hours saved is convincing for an operations head; capital or regulatory-risk reduction, for a committee. The same project is argued differently depending on who signs, without either version ceasing to be true.

- **Prefer a supported range to an exact figure with nothing behind it.** *"Between 1.5 and 2.5 full-time equivalents, depending on the scope of phase two"* is more credible and more defensible than a point estimate nobody can sustain.
- **Quantify the cost, not just the benefit.** Include what is usually left out: the time of a team that is already busy, second-year maintenance, and the learning curve.
- **Name the main risk yourself.** If it is not in the document, it will come up in the forum and on less favourable terms. Stating it alongside its mitigation keeps control of the framing.
- **Answer "why now".** A business case is rarely rejected for being poor; more often it is deferred. If there is no reason for this quarter rather than the next, deferral is the likely outcome.

18

Difficult conversation downward: asking someone on your team to change a behaviour

When one person's conduct affects the team, the cost of not having the conversation is absorbed by the rest of the team, and absorbed in silence. Postponing it looks like consideration, but it transfers the problem to those who did not cause it. The next five sections cover the preparation, the distinction that decides the outcome, the script for the conversation, the predictable reactions, and what is worth documenting afterwards.

18.1 Preparation: three columns before speaking

Before the conversation it is worth having three things in writing:

- **An observable fact.** What happened, when, and in what context. Verifiable, specific and without adjectives: something a third party present would describe the same way.
- **The impact.** What consequence it had on the work or on the team, not on the feelings of whoever raises it. *"Two people stopped bringing items to the session"* is impact; *"it bothers me"* is not.
- **The expected change.** What has to be different, specifically, and what success would look like in four weeks. If that is hard to describe, it is worth observing a little longer before having the conversation.

If filling them in produces impressions rather than facts, the problem may still be real, but it is worth observing first. A conversation built on perceptions turns into an argument about whether the perception is fair, and that argument has no exit.

18.2 Behaviour, not identity (*observable behaviour, not traits*)

This distinction decides the outcome of the conversation. An observation about behaviour admits correction; an observation about identity admits only defence, because it tells the person who they are.

Identity (closes down)	Behaviour (opens up)
You are arrogant with the team.	In the last three sessions you interrupted the analysts before they finished presenting.
You are not collaborative.	You closed the last two shared deliverables without telling the two people who depended on them.
You have a bad attitude.	When you were asked to adjust the report, you replied in the team chat that it was a waste of time.

18.3 The script, in five beats

1. **A direct opening**, without preamble and without preliminary praise: *"I want to talk about something that is affecting how the team works."* Praise beforehand makes the person brace for criticism and devalues any praise that follows later.
2. **The fact.** One or two, the clearest, not the full list: a long list is received as a file and pushes the conversation towards defence.
3. **The impact**, in concrete terms of work.
4. **Effective listening.** Put the question and hold the pause. This is where the method meets coaching practice. Information may surface that changes your reading of the matter, and if it does, it is worth taking in properly rather than appearing to.
5. **A verifiable agreement and a follow-up.** What changes, how it will show, and when it gets reviewed. Without a review date, the agreement rarely holds.

18.4 The three reactions and what to do

Three reactions cover most cases:

- **Denies the fact.** Return to the fact, not the judgement. If the disagreement is about what happened, that is the point to settle before any other.
- **Reacts with distress.** Slow the pace and hold the message. Withdrawing it out of discomfort is the most frequent error and forces the conversation to be repeated in worse conditions. It is possible to be firm and considerate at once.
- **Responds with counter-charges.** Acknowledge whatever is legitimate, schedule it separately with a date, and return to the original subject. Handling both at once usually leaves both unresolved.

18.5 What to document afterwards

A short email after the conversation — what was agreed and the review date — protects both parties and prevents a later argument about what was said. It is not a file and needs no formal tone; it is shared memory. If the situation later escalates, that record is what distinguishes an orderly process from one person's word against another's.

19

Difficult conversation upward: when the problematic behaviour is your manager's

This is a different case, and it is best not treated as the previous one with the roles reversed. Downward, a change is asked for with authority; upward, that authority does not exist and the cost of error is asymmetric. That does not mean staying silent, but preparing the conversation more carefully. The next three sections cover what changes relative to the previous case, how the conversation is structured, and how to calibrate the risk before having it.

19.1 What changes relative to the previous case

Four differences:

- **A change of behaviour is not what is asked for**; an effect is set out and a bounded request is made. The conclusion still goes first, but the conclusion is the request, not a diagnosis of their character.
- **The impact is measured in results, not in climate.** *"When the priority changes mid-week without notice, two deliverables slip"* admits discussion. *"The team feels demotivated"* invites them to question your reading of the team.
- **The conversation is asked for.** One sentence is enough and it changes the whole tone: *"Can I raise something about how we are working? It is awkward, but I think it is worth your knowing."*
- **A way out is offered.** A concrete proposal that makes it easier to do differently, not just the observation. Without a proposal, the conversation reduces to a well-drafted complaint.

19.2 Proposed structure

> *"Do you have ten minutes? I would like to propose an adjustment to how priority changes reach the team.*
>
> *Over the last three weeks, two of the committed deliverables slipped because the focus changed mid-week. The team managed to cover the urgent items, but we were late with finance on two occasions.*
>
> *My proposal is to agree priority changes on Mondays, and for anything arriving later to wait unless it is regulatory. If something does come in outside that window, I will tell you which deliverable moves, you decide, and we are both covered.*
>
> *Shall we try it for a month?"*

It is worth noting what the framing does not contain: no adjectives about the other person, no generalisation and no accumulation of grievances. A measured effect, a proposal and a trial period. That is hard to argue with, and if it is rejected anyway, it yields valuable information about the real room available.

19.3 Calibrating the risk

Four considerations before having it:

- **Channel and timing.** Never in public, never immediately after the episode, never in writing as the first contact. Email closes options because it is on the record; a conversation leaves room for both sides to adjust.
- **One thing per conversation.** Using the moment to raise everything that has accumulated is what turns a reasonable request into a complaint.
- **When it is not worth it:** when the pattern is structural and sustained, when this method has already been tried without result, or when the personal cost exceeds the benefit that can be expected. Recognising that is not cowardice; it is calculation. Most of these conversations go better than anticipated; the highest cost is usually the year spent avoiding them.
- **Escalating or going to human resources** carries a real and often irreversible cost to the relationship. It is the right route for conduct that crosses an ethical or behavioural line, and an expensive route for management-style problems. Telling which is which, before acting, is part of the work.

20

Structuring a "no" (saying no)

Saying no badly is one of the most persistent sources of overload. A no wrapped in three paragraphs of context reads as *"let me look into it"* and comes back the following week as a commitment. The next four sections cover the rule, the three kinds of no, how it is structured, and an example of each.

20.1 The rule

The no goes first, in the first sentence, without preamble. Then the reason, then the alternative. That order is more considerate than the long version: it lets the other party stop reading and look for another route straight away, instead of waiting for an answer that is already settled.

20.2 Three kinds of no, and the error of confusing them

- **A capacity no.** It can be done, but something else moves. Better not stated as *"I can't"*, but as an explicit trade: *"Yes, if we move the X delivery by two weeks."* That returns the decision to whoever has the authority to prioritise, which is where it belongs.
- **A scope no.** It does not belong to your team, or it belongs to another that will handle it better. Here it matters to name who does, because a no with no destination reads as evasion.
- **A judgement no.** It can be done and should not be. It is the most uncomfortable and the one that most defines you professionally. It requires an explicit technical reason and a willingness to hold it. It is also the only one that admits no softened alternative: offering a reduced version of what should not be done amounts to abandoning your own argument.

20.3 Proposed structure

1. **The no, explicit, in the first line.**
2. **One reason**, the real and verifiable one. Just one: three reasons read as accumulated excuses, and refuting one is enough to reopen the matter.
3. **The alternative, if there is one:** another route, another owner, another date, or a smaller scope.
4. **What would make it a yes.** This is what distinguishes a collaborative no from a closed refusal, and it often unlocks the underlying conversation.

20.4 One example of each kind

A capacity no:

> *"Thank you for the proposal. We will not manage to include the sensitivity analysis in Thursday's version: the team closes the regulatory exercise that same week and there is no room without putting that date at risk.*

> *Two options: we deliver it complete on the 12th, or Thursday's version goes without sensitivity and the annex follows. My recommendation is the second, because the committee can decide without that analysis.*
>
> *If the sensitivity is essential for Thursday, we would need to move the regulatory exercise, and that would require your approval. Happy to discuss."*

A scope no:

> *"Thank you for considering the team for the liquidity model review. It is not an analysis that belongs to this function: the liquidity methodology is maintained by Treasury, which also already holds the history.*
>
> *If it helps, I can put you in touch with the person who owns it this week, and share the context so you do not have to repeat it.*
>
> *Happy to help with anything else."*

A judgement no:

> *"Thank you for the confidence, but I would not recommend publishing the indicator on the current methodology.*
>
> *The calculation does not distinguish between the two segments, so the published value would not be comparable with last quarter's and could support mistaken conclusions.*
>
> *What we can do is publish the series once the segments are separated, which we estimate at three weeks.*
>
> *If the date cannot move, the indicator would have to go out with an explicit methodological note, and in that case I would ask for the decision to be documented. Happy to discuss."*

The judgement example does not offer a reduced version of what is being objected to — that would contradict the previous section — but a different route. And where the date cannot move, it asks for the decision to be documented, which is the honest exit when preventing something you disagree with is not in your hands.

21

Handing over a project or a system (handover, knowledge transfer)

Almost all handovers are documented badly for the same reason: whoever is handing over documents what they know, not what the recipient will need. And what the recipient will need is, above all, what the person handing over no longer notices because it has become second nature. The next four sections cover when a handover is considered finished, what the document must contain, in what sequence the transfer happens, and how long it should be.

21.1 The handover ends when the recipient completes a cycle unaided

A handover does not end with the delivery of documentation. It ends when the recipient runs a complete cycle without support and the result is correct. Defining it that way changes everything else: the central deliverable is not a document but a supported sequence with a cut-off date.

21.2 What the document must contain

Seven sections:

1. **What it is and why it exists**, in one paragraph. Who uses it and what decision depends on it.
2. **Current state, with nothing left out.** What works, what is half-done, and what runs on a temporary workaround. Omitting technical debt does not remove it: it guarantees the recipient finds it at the worst moment and associates it with the previous tenure.
3. **How it is operated.** Routines, calendar, inputs, dependencies, who delivers what and when, with concrete names and dates.
4. **Who is who.** Not just contacts: what each one cares about and where there is history. It is the section least often written and the one most needed.
5. **Tacit assumptions.** The three to five *"this is always done this way because..."* that appear in no document. It is the most valuable section and the one that takes most effort to recover.
6. **Known risks and pending decisions**, with a recommendation on each.
7. **What to do when it fails.** The three or four frequent failure modes and their fix, written as executable instructions.

21.3 The transfer sequence

Four stages, each with at least one real cycle. None is skipped even when the recipient is experienced, because what is being transferred is not technical capability but context.

1. **The person handing over runs it; the recipient observes.**

2. **The recipient runs it; the person handing over observes** and does not intervene except on a material error.
3. **The recipient runs it** with the person handing over available but not present.
4. **The recipient runs it autonomously.** A declared cut-off date and a short warranty period — two or three weeks — in which the person who handed over answers questions but does not execute.

> **The cut-off date has to exist and be public. Without it, handovers become permanent: the recipient keeps asking, the person who handed over keeps solving, and six months later the system is still theirs in practice and the recipient's on the org chart.**

Record the first run. Stage 1 is the natural moment to record the whole process on video, narrating what is being done. A recording handles well what a document describes badly — which screen, in what order, what it looks like when the result is correct — and lets the written document stay short. Three conditions for it to be useful: that it lasts as long as the process and no longer, that it is stored where the recipient can find it without having to ask, and that it is re-recorded when the process changes substantively.

In a regulated environment one limit is worth checking before recording: if the screen shows client data or information subject to confidentiality, recording may be restricted or require masking. Confirming that with the relevant function beforehand, rather than afterwards, avoids having to destroy the material.

21.4 A short document that is kept current beats a long manual

An eighty-page manual is rarely consulted. What gets used is a short operating document — five to ten pages — the technical files wherever they live, and, where it exists, the recording of a complete cycle. If time is short, it is best to write the tacit-assumptions section and the frequent-failures section first: they are the only ones nobody else can write.

22

Explaining a technical subject in plain terms (pitching to a non-technical audience)

The underlying misunderstanding is believing that simplifying means giving up rigour. It does not: simplifying well consists of changing the level of abstraction without saying anything false, and it is harder than explaining in detail, because it requires determining which part of the detail is dispensable for that particular audience. The next four sections cover the layers the explanation is organised into, the rules that hold it together, how it looks applied, and how to verify it was understood.

22.1 The three layers

1. **What this decides.** The practical consequence for the listener. It goes first and is often all that is needed.
2. **How it works, in one sentence.** The mechanism at its most general level, with an analogy if that helps.
3. **The detail, available but not imposed.** Better offered than pre-empted: *"If you are interested, I can show you how it is calculated"* works better than working through the calculation aloud when nobody asked.

22.2 Practical rules

Six rules:

- **Start with the consequence.** *"This change means we approve fewer customers in the mid segment"* before any methodological consideration.
- **Use an analogy and say where it stops holding.** Every analogy fails somewhere; flagging that keeps the rigour and stops anyone building a mistaken conclusion on it.
- **Replace jargon with the effect, not with a synonym.** Instead of explaining what calibration is: *"the model ranks customers well from best to worst, but the absolute level of risk it assigns is too low"*.
- **Anchor numbers to something known.** *"Four in every thousand loans"* communicates more than *"0.4%"*, and *"the equivalent of two weeks of the team"* more than *"80 hours"*.
- **Avoid the phrase "it's complicated".** It is received as condescension and as evasion. If the subject does not fit the time available, it is better to say which part does.
- **Avoid defensive detail.** A long technical explanation often does not aim at the other person's understanding but at demonstrating command of the subject. The effect is the opposite: it subtracts credibility rather than adding it.

22.3 Worked example

Technical version, aimed at the wrong audience:

> BEFORE
>
> *"The recalibration adjusts the intercept of the link function to align the long-run average probability of default with the observed through-the-cycle rate, without modifying the coefficients or the model's discriminatory power."*

The same information, for a business audience:

> AFTER
>
> *"The model still ranks customers with the same accuracy: the ordering from best to worst does not change. What is wrong is the level. The model assigns 2% default where observed experience over the last three years has been 3%, and we are going to correct it to 3%.*
>
> *In practice, some customers who pass the cut-off today will stop passing it, mostly in the mid segment. I have the estimate of how many and in what proportion.*
>
> *What does not change is the ranking, so relative commercial decisions — who gets offered a larger limit — stay as they are."*

22.4 Ask them to explain it back to you

Ask the person to explain it back, naturally: *"how would you tell your team?"*. What comes back is exactly what they understood. It is uncomfortable the first time and it is the only reliable diagnostic available: nodding, or not asking, does not indicate comprehension.

PART III

Practice

How to install the habit, in yourself and in a team.

23

Common failures in a quantitative profile

Ten patterns that recur in technically sound documents. Each points back to the Part I section that covers it:

1. Reproducing the order of the analysis instead of the order of the decision.
2. Confusing precision with relevance: putting the confidence interval in the executive summary and leaving out the recommendation.
3. Topic headings instead of message headings.
4. Lists of findings with no summary idea: the structure exists, but the message does not land.
5. Long deductive chains at the top level, forcing the reader through everything before understanding anything.
6. Grouping by data convenience — whatever came out of the same query — instead of by the logic of the argument.
7. Annexes carrying the conclusion: the key message hidden in table 7 of appendix B.
8. Avoiding the conclusion out of political caution, with the result that the document contains nothing verifiable.
9. Assuming the reader knows the jargon. Every acronym is spelled out the first time, including the obvious ones in your own function.
10. Fixing the prose when the problem is the analysis. If "so what?" produces no answer, no rewrite resolves it.

24

Checklist

Worth using before sending any document or presentation. Ten questions, five minutes.

1. Is the first substantive sentence my conclusion, rather than an announcement that one is coming?
2. Can I state the situation and the complication in two lines each?
3. Is the question my introduction generates the one my document answers?
4. Does every group have between two and five items of the same kind?
5. Does every summary idea survive the "so what?" test?
6. Does every group follow a time, structural or degree order, and can I name which?
7. Are the groups mutually exclusive, and what would the best-informed sceptic say is missing?
8. Reading only the headings, does the full argument come through?
9. Is the executive summary enough to decide without reading the rest?
10. Is every acronym spelled out the first time it appears?

25

A 30-day practice plan

Pyramid structure is not installed by reading. These four weeks are designed around material you have to produce anyway, with no meaningful extra time.

Week	Exercise	Time
1	Take three documents you have already sent. For each, write the SCQA introduction it should have had and the conclusion in one sentence. Do not rewrite them in full.	20 min each
2	Every email longer than five lines: request or conclusion in the first line, no exceptions. Before sending, apply "so what?" to that first line.	0 extra min
3	Take your last presentation. Rewrite every title as a full sentence with a conclusion. Read them aloud in order and find where the argument breaks.	1 hour
4	Write a new document entirely with the method, top-down. Before sending it, give a colleague only the executive summary and ask what decision is being requested.	2 hours

From the second month, the only habit to sustain is the "so what?" test applied to every heading and every slide title. Everything else becomes automatic.

26

Running this as a team workshop

Handing out the document does not teach the method. What works is a short session in which each participant rewrites material of their own and receives criticism on it. The ninety-minute agenda that follows splits the time in that proportion: eighty per cent is taken up by participants writing, and not by whoever runs the session talking. The next four sections cover the preparation beforehand, the agenda, the rules for the criticism, and what to expect from the result.

26.1 Before the session

Three conditions beforehand:

- Ask each participant for a document of their own that has already been sent: a memo, a long email, a summary slide. Their own material, real, not an example built for the exercise; without that the session does not work.
- Send the guide in advance, without assuming anyone will read it. The session has to work without prior reading.
- Eight people maximum. Beyond that, nobody receives individual criticism and the exercise dilutes.

26.2 Agenda

Block	Content	Minutes
Opening	An example of a badly structured document — one of your own from some time ago, or the one in section 14 of this guide — and the question to the group: what is being asked here? It is worth letting them work at it unaided: that bewilderment is the argument for the workshop.	10
Method	Only three things: conclusion first, SCQA, and the "so what?" test. MECE and the deduction–induction distinction are better left to a later session.	20
Exercise 1	Each participant writes the conclusion of their document in one sentence, and applies "so what?" until it produces nothing new.	15
Criticism	Three volunteers read their sentence. The group asks "so what?" out loud. It is uncomfortable, and it	20

	is where the learning happens.	
Exercise 2	Each participant rewrites the introduction of their document using SCQA, in four lines.	15
Close	The concrete commitment: for the next month, every email longer than five lines carries the request in the first line. Nothing else.	10

26.3 Rules for the criticism

Four rules:

1. The document is criticised, never the person. It is the same behaviour-versus-identity distinction from section 18, applied to a text.
2. Whoever criticises proposes a concrete alternative. "It is not clear" is not useful criticism; "I would put this sentence first" is.
3. The author does not defend during the criticism: take notes and decide later. Defending in the moment turns the exercise into a debate and consumes everyone's time.
4. Whoever runs the session goes first, with a document of their own. Without that, nobody else really exposes themselves.

26.4 What to expect

The first session usually reveals something more uncomfortable than a writing problem: that several documents have no conclusion because none was ever reached. That is not a failure of the workshop, but its most valuable outcome. Worth saying so out loud at the start, so that it is not experienced as exposure.

The habit does not install itself in one session. What sustains the change is what happens afterwards: if documents from the team are consistently returned with the same question — what is the conclusion, and does it come first? — in two months the team writes differently. Without that follow-up, the session is remembered, but the habit does not change.

27

What this method does not solve

It is worth being explicit about the limits, because the method is sometimes presented as a general solution to communication and it is not. Four caveats:

- It does not improve spoken presence. Structuring an argument and holding it under interruption, pressure or hostile questions are different skills. Those are trained by repetition in front of an audience with feedback, not by writing.
- It does not substitute for analysis. If "so what?" produces no answer, what is missing is analytical work, not structure.
- It does not resolve organisational politics. An impeccable document addressed to the wrong person still fails to produce the intended effect.
- It does not generate business intuition. It helps you express the judgement you already have; it does not create it.

With those four caveats in mind, it remains the method with the highest return per hour invested for someone technically solid whose bottleneck is that the content does not reach the decision-maker in order.

28

Glossary

These are the terms as they are used in the method and in day-to-day work in global organisations, with a short definition of what each one names.

Acronyms

Acronym	Stands for	What it is
SCQA	Situation, Complication, Question, Answer	The introduction pattern.
MECE	Mutually Exclusive, Collectively Exhaustive	Quality standard for a grouping: no overlaps, no gaps.
BLUF	Bottom Line Up Front	Writing convention: the conclusion in the first line.
SCR	Situation, Complication, Resolution	A three-part variant of the same introduction pattern.

Terms of the method

Term	What it names
Governing thought / main message	The single statement that governs the whole document — the conclusion that leads the pyramid.
Summary statement	The sentence that summarises a group of ideas and moves up a level.
Vertical relationship	Each idea answers the question raised by the one above.
Horizontal relationship	Ideas at the same level are of the same logical kind.
Top-down	Building from the conclusion, when you already have it.
Bottom-up	Building by grouping loose findings until you reach the conclusion.
Time order	The items are steps or linked causes.
Structural order	The items are parts of a whole.
Degree order / ranking order	The items are ordered by importance or magnitude.
Intellectually blank assertion	A sentence that announces content without delivering it.
So what? test	Asking until the idea carries actionable content.
Disturbing event	What broke the stable situation and motivates the document.
Undesired result (R1) / desired result (R2)	The two ends that define the problem.

Diagnostic framework	A decomposition used to find the cause.
Logic tree	A decomposition of options with explicit criteria.
Message title	A heading that states the conclusion, not the topic.
Action title	A slide title written as a sentence with a conclusion.
Storyline	The argument that reads straight through from the titles alone.
Executive summary	The conclusion plus the first level of support, in prose.

Terms of the working environment

Term	What it names
Data owner	Whoever is formally accountable for the quality of a source.
Escalation path	The agreed sequence for raising an issue without damaging the relationship.
Business case	The document used to request budget.
Base case / do-nothing option	What happens if nothing is approved. The real comparison.
Opportunity cost	What stops being done if this is approved.
Stage gates	Splitting the investment into tranches with a review between them.
Pre-wire / socialise	Talking the matter through with each decision-maker before the committee.
Policy exception / waiver	Formal permission to depart from a standard.
Intent vs impact	Distinguishing what someone meant from what they produced.
Feedback sandwich	Praise, criticism, praise. The recipient retains the praise and not the message.
Handover / knowledge transfer	Transferring a system or project to another person.
Shadowing	The stage in which the receiver observes before executing.
Runbook	The document with routine and incident procedures.
Technical debt	Temporary fixes that are still in place and accumulating cost.
Curse of knowledge	Taking as obvious what took years to understand.
Order of magnitude	The scale of the number, before the exact number.

29

Sources and attribution

This document is original material: the text, the examples and the cases were written for this guide and do not reproduce content from any of the works cited. What it does take from them are the methods, which belong to their authors and are credited here.

The core method

The pyramid structure, the SCQA pattern, the MECE standard, the distinction between time, structural and degree order, and the problem-definition framework all come from The Pyramid Principle: Logic in Writing and Thinking, by Barbara Minto (Pearson / Financial Times). Anyone who wants the original formulation and its examples should go to the book; this guide does not replace it.

Other sources

Idea	Origin
The behaviour-versus-identity and intent-versus-impact distinctions, in sections 18 and 19	Difficult Conversations: How to Discuss What Matters Most, by Douglas Stone, Bruce Patton and Sheila Heen (Harvard Negotiation Project).
Diagnosis as the first component of a decision, in section 17	Good Strategy Bad Strategy, by Richard Rumelt.
The curse of knowledge, in section 22	A behavioural economics concept coined by Camerer, Loewenstein and Weber, popularised by Chip and Dan Heath in Made to Stick.
The habit of asking before solving, in the approach to conversations	The Coaching Habit, by Michael Bungay Stanier.

About the examples

All examples are invented. The figures, portfolios, models, borrowers and situations described do not correspond to any specific institution, and any resemblance to a real case follows from the fact that these problems recur everywhere.

This guide is shared non-commercially for team use. If you find it useful, the best gesture towards the authors cited is to buy their books.